

# Information Retrieval Assignment Report

Assignment 3: Text Ranking with Transformer-Based Models

Naveeta Maheshwari (Roll No: 2024AIZ8309)

Sagar Singh (Roll No: 2025SIY7574)

**Course:** Information Retrieval and Web Search (COL764/COL7364)

**Semester:** 2025–26

November 28, 2025

## 1 Introduction

The retrieve-then-rerank paradigm is a standard approach to efficient and accurate ranking. A lexical retriever (BM25) returns a set of candidates quickly; a heavy neural rerank (Hugging Face cross-encoder) computes relevance scores over  $\langle q, d \rangle$  pairs and reorders the candidates. The final reported list is truncated to the top-10 for evaluation.

## 2 Task 1: Implement the ”Retrieve and Rerank” framework using BERT

### 2.1 Implementation Details

#### 2.1.1 Index and Retriever

We used **Pyserini**’s prebuilt `msmarco-passage` Lucene index for BM25 retrieval. For each query, we retrieve top- $k$  documents with  $k \in \{10, 25, 50, 75, 100, 150\}$ .

#### 2.1.2 Reranker: MiniLM Cross-Encoder

We use the fine-tuned model `cross-encoder/ms-Marco-MiniLM-L-6-v2` from HuggingFace as the heavy reranker.

MiniLM is a distilled version of BERT trained with knowledge distillation from large transformer models, reducing number of layers (6 vs. 12 in BERT-base) while preserving strong contextual relevance matching. This makes it significantly more efficient for inference in retrieval pipelines. As a **cross-encoder**, the model jointly encodes the query–document pair in a single input sequence:

[CLS]  $q$  [SEP]  $d$  [SEP]

Unlike bi-encoders that independently embed  $q$  and  $d$ , the cross-encoder allows **full cross-attention** between every query and document token. This enables the model to learn rich, token-level semantic interactions crucial for fine-grained relevance estimation in MS MARCO.

### 2.1.3 relevance score for reranking

The final [CLS] embedding is passed through a classification layer to produce a real-valued **relevance score**:

$$s(q, d_i) = f_{\text{MiniLM}}([\text{CLS}])$$

We apply this scoring to each of the top- $k$  BM25 candidates independently, sort by score, and then output the final top-10 for evaluation.

Since the model must score each  $\langle q, d \rangle$  pair separately, complexity grows linearly as  $\mathcal{O}(k)$ . This accounts for the approximately linear increase in total latency observed in Figure ??, while yielding consistent gains in NDCG@10 and MRR@10 as  $k$  grows.

### 2.1.4 Chunking

MS MARCO passages are short; in our runs, input lengths remained within the cross-encoder’s maximum sequence length. Therefore, no chunking or score aggregation was required.

**Outputs and Evaluation.** For each  $k$ , we produce two TREC-style run files: `bm25_k{K}.run` and `rerank_k{K}.run` with lines of the form: `qid docid rank score`. We compute **NDCG@1/5/10** and **MRR@10** using our evaluation script (qrels format: `qid docid rel`). We also log average per-query latency for BM25, reranking, and total.

## 2.2 Results and Observations

### 2.2.1 NDCG

**BM25:** The BM25 performance remains flat across all  $k$  values, as its top-ranked set does not change with a larger retrieval pool.

**Reranker:** The reranker consistently improves with increasing  $k$ , showing higher NDCG scores that stabilize around  $k = 100$ – $150$ , demonstrating its ability to identify and prioritize truly relevant documents. The corresponding figures are shown in Figure 1.




Figure 1: NDCG scores vs.  $k$  for BM25 and Rerank.

### 2.2.2 MRR@10

**BM25:** MRR@10 remains nearly constant at 0.534 across all  $k$  values.

**Reranker:** MRR@10 rises from approximately 0.607 at  $k = 10$  to around 0.627 at  $k = 100$ , then fluctuates slightly, reflecting better early precision as relevant documents

are promoted.

The corresponding plot is shown in Fig. 2.




Figure 2: MRR@10 vs.  $k$  for BM25 and Rerank.

### 2.2.3 Best Results for Task1

Table 1: Rerank evaluation results at different  $k$  values. Best results per metric are underlined.

$k$	NDCG@1	NDCG@5	NDCG@10	MRR@10
10	0.403	0.327	0.297	0.607
25	0.363	0.342	0.330	0.582
50	0.394	0.345	0.349	0.605
75	0.394	0.350	0.353	0.604
100	<u>0.403</u>	<u>0.362</u>	<u>0.367</u>	<u>0.627</u>
150	<u>0.403</u>	0.357	<u>0.373</u>	0.622

## 2.3 Conclusion

The Retrieve & Rerank framework with a fine-tuned BERT cross-encoder consistently outperforms BM25 across all evaluated metrics, with the largest gains observed at higher  $k$  values. While reranking effectiveness improves as  $k$  increases—reflected in rising NDCG@1/5/10 and MRR@10—marginal returns diminish beyond  $k \approx 100$ , and latency grows linearly. Therefore, choosing an appropriate  $k$  balances improved ranking quality against computational cost, depending on application latency constraints. Overall, our framework demonstrates that augmenting BM25 with cross-encoder reranking delivers significant precision gains with modest additional overhead.

### 3 Task 2: Comparisons with BERT alternatives

This task investigates two additional reranking baselines on the MS-MARCO dataset:

1. ELECTRA cross-encoder
2. MonoT5 generative reranker

We evaluate their performance across different candidate pool sizes  $k \in \{10, 25, 50, 75, 100, 150\}$  using NDCG@1, NDCG@5, NDCG@10, and MRR@10..

#### 3.1 Model Architecture Comparison

Table 2: Comparison of retrieval and reranking models with differences from BERT.

Model	Architecture	Difference from BERT	Key Characteristics
BM25	Lexical inverted index	Not a neural model; purely term-based	Fast retrieval; no semantic understanding; retrieves candidate pool.
ELECTRA Cross-Encoder (baseline-1)	Transformer encoder fine-tuned for relevance	Uses ELECTRA discriminator pre-training instead of BERT’s masked LM; cross-encodes query-document pairs	Strong token interaction modeling; performance improves as $k$ grows; cost $\mathcal{O}(k)$ .
MonoT5-Small (baseline-2)	Encoder-decoder (Seq2Seq generative transformer)	Seq2Seq generative model vs. BERT’s encoder-only architecture	Weaker discrimination for short passages; performance <i>degrades</i> as $k$ grows due to noisy candidates.

#### 3.2 Results and Observation

##### 3.2.1 Results for MonoT5

Table 3: MonoT5 Baseline performance across candidate set size  $k$ . Best results per metric are underlined.

$k$	NDCG@1	NDCG@5	NDCG@10	MRR@10
10	<u>0.2952</u>	<u>0.2728</u>	<u>0.2635</u>	<u>0.5085</u>
25	0.2724	0.2610	0.2695	0.4837
50	0.2610	0.2315	0.2682	0.4926
75	0.2438	0.2166	0.2569	0.4804
100	0.2238	0.2028	0.2474	0.4701
150	0.2010	0.2064	0.2380	0.4562

### 3.2.2 Results for ELECTRA Cross-Encoder

Table 4: ELECTRA Cross-Encoder performance across candidate set size  $k$ . Best results per metric are underlined.

$k$	NDCG@1	NDCG@5	NDCG@10	MRR@10
10	0.3371	0.3179	0.2908	0.5545
25	0.3771	0.3308	0.3165	0.5760
50	0.4314	0.3646	0.3522	0.6379
75	0.4400	0.3773	0.3615	0.6335
100	<u>0.4800</u>	<u>0.3996</u>	0.3856	<u>0.6710</u>
150	0.4429	0.3914	<u>0.3918</u>	0.6653

### 3.3 Plots: MonoT5 Trends




Figure 3: MonoT5 performance metrics across candidate set size  $k$ .

### 3.4 ELECTRA Cross-Encoder: Effectiveness and Latency

plots for ELECTRA Cross-Encode are give in Fig. 4




Figure 4: Effectiveness of ELECTRA reranker across candidate pool sizes  $k$ .




Figure 5: Latency vs  $k$  for ELECTRA reranker (BM25, Rerank, Total).

### 3.5 Analysis & Observations

- **MonoT5 performance drops as  $k$  increases:** Expanding the candidate pool introduces more irrelevant passages, which the pointwise T5 classifier struggles to discriminate.

- **ELECTRA cross-encoder improves steadily with  $k$ :** Cross-attention allows better exploitation of larger  $k$  because relevant items are more likely to enter the pool.
- **BM25-only remains flat:** Final top-10 ranking unchanged across  $k$ .
- **Takeaway:** ELECTRA achieves a better accuracy–cost trade-off. MonoT5 is slower and degrades with noise.
  - The ELECTRA Cross-Encoder achieves its best NDCG@1 (0.4800), NDCG@5 (0.3996), and MRR@10 (0.6710) at  $k = 100$ , while the highest NDCG@10 (0.3918) occurs at  $k = 150$ .

**Conclusion.** Task-2 shows that although MonoT5 is a widely-used generative reranker, it is *inferior* to ELECTRA-based cross-encoding for short MS MARCO passages. For production usage, a cross-encoder like MiniLM/ELECTRA appears to be the correct compromise between performance and latency.

## 4 Task 3: Improving the Best Reranker (Pseudo-Relevance Feedback)

From Tasks 1 and 2, the **ELECTRA cross-encoder** achieved the best ranking performance. In this task, we attempt to further improve its effectiveness by introducing **Pseudo-Relevance Feedback (PRF)** for query expansion prior to reranking.

### 4.1 Method

For each query:

1. Retrieve top- $k$  documents using BM25.
2. Select the **top-5 documents** as pseudo-relevant.
3. Extract the **top-5 most frequent terms** from those documents.
4. Expand the query by appending these feedback terms.
5. Re-rank the expanded query’s candidate set using:

`cross-encoder/ms-marco-electra-base`

This method assumes the highest-ranked BM25 documents are relevant and can provide useful context to enhance the neural reranker’s semantic matching ability.

### 4.2 Results

We evaluate the modified pipeline across  $k \in \{10, 25, 50, 75, 100, 150\}$  using the same metrics as in Tasks 1 and 2 (NDCG@1/5/10 and MRR@10).

**Finding.**

- At  $k = 10$ , both **MRR@10 and NDCG@1 showed improvement**, indicating that when BM25 returns highly relevant candidates, feedback terms strengthen the query.
- For  $k \geq 25$ , performance **degraded consistently across all metrics**. The expanded query introduced noise from non-relevant documents, hurting the neural reranker’s discrimination ability.

### 4.3 Analysis

PRF behaves differently in dense reranking environments than in classical IR:

- Neural cross-encoders already capture fine-grained relevance.
- Naively adding frequent lexical terms may:
  - reduce semantic precision,
  - introduce unrelated context,
  - confuse relevance modeling.
- PRF only helps when the initial ranking quality is already high (as seen at  $k = 10$ ).

## 4.4 Conclusion

While pseudo-relevance feedback is traditionally effective in retrieval, it is **not directly beneficial in neural reranking** unless the feedback set is clean.

Overall, our experiments demonstrate that integrating classical PRF techniques into cross-encoder reranking requires careful control of feedback noise to avoid degrading performance.